

Random Forest and Boosting

Default Modeling

Contents

Random Forest and Boosting	2
1. Upload and prepare data	2
1.1 Upload data	2
Create default_event	3
Create default flag	3
1.2 Select a subset of variables	3
1.3 Filter out NAs	3
1.4 Split train/test	3
Perform Random Forest analysis	3
2.1 Fit random forest	3
2.2 Variable importance analysis	4
3. Perform boosting analysis	6
3.1 Test sample analysis	8
3.2 Inclusion of shrinkage	8
Return the Optimal Number of Iterations	10
ML Calibration	12
Create the data set and fit calibration function	12
Fit the calibration function	12
ML Model validation	13
Calibrated PD Validation	14

```
library(dplyr)
library(caret)
library(randomForest)
library(gbm)
library(ROCR)
library(optiRum)
library(smbinning)
```

Random Forest and Boosting

1. Upload and prepare data

1.1 Upload data

```
oneypd_tree <- read.csv(file = 'Z:/Model Risk/Adam/IFRS9_CECL_MV/data/chap2oneypd.csv')
dplyr::glimpse(oneypd_tree)
```

```
## Observations: 25,906
## Variables: 45
## $ X               <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 1...
## $ id              <int> 6670001, 9131199, 4963167, 39185...
## $ vintage_year     <int> 2005, 2006, 2004, 2005, 2006, 20...
## $ monthly_installment <dbl> 746.70, 887.40, 1008.50, 458.23,...
## $ loan_balance     <dbl> 131304.44, 115486.51, 128381.73,...
## $ bureau_score     <int> 541, 441, 282, 461, 466, 470, 51...
## $ num_bankrupt_iva <int> 0, 0, 0, 0, 0, 0, 0, 0, NA, 0, 0...
## $ time_since_bankrupt <int> 0, 0, 0, 0, 0, 0, 0, 0, NA, 0, 0...
## $ num_ccj          <int> 0, 0, 1, 0, 0, 0, 0, 0, NA, 0, 0...
## $ time_since_ccj    <int> 0, 0, 36, 0, 0, 0, 0, 0, NA, 0, ...
## $ ccj_amount        <int> 0, 0, 459, 0, 0, 0, 0, 0, NA, 0,...
## $ num_bankrupt      <int> 0, 0, 0, 0, 0, 0, 0, 0, NA, 0, 0...
## $ num_iva           <int> 0, 0, 0, 0, 0, 0, 0, 0, NA, 0, 0...
## $ min_months_since_bankrupt <int> 0, 0, 0, 0, 0, 0, 0, 0, NA, 0, 0...
## $ pl_flag           <int> 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0,...
## $ region            <fct> r_a, r_b, r_c, r_d, r_e, r_c, r...
## $ ltv               <dbl> 0.7586, 0.6973, 0.6959, 0.1099, ...
## $ arrear_months     <dbl> 0.0000000, 0.0000000, 2.1882300,...
## $ origination_date  <fct> 9/14/2005, 1/20/2006, 12/21/2004,...
## $ maturity_date     <fct> 9/30/2040, 1/31/2031, 12/31/2029,...
## $ repayment_type    <fct> Non-IO, Non-IO, Non-IO, Non-IO, ...
## $ arrear_status     <int> 1, 1, 2, 1, 1, 1, 1, 1, 2, 1, 2,...
## $ arrear_segment    <int> 1, 1, 2, 1, 1, 1, 1, 1, 2, 1, 1,...
## $ mob               <int> 120, 116, 129, 123, 110, 120, 13...
## $ remaining_mat     <int> 300, 184, 171, 93, 310, 0, 166, ...
## $ loan_term         <int> 35, 25, 25, 18, 35, 10, 25, 25, ...
## $ live_status       <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,...
## $ repaid_status     <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,...
## $ month             <int> 9, 9, 9, 9, 9, 9, 9, 9, 9, 9, 9,...
## $ arrear_event      <int> 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0,...
## $ bankrupt_event    <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,...
## $ term_expiry_event <int> 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,...
## $ worst_arrear_status <int> 1, 1, 2, 1, 1, 1, 1, 1, 2, 1, 2,...
## $ max_arrear_12m    <dbl> 0.0000000, 0.0000000, 2.188230, 0....
## $ recent_arrear_date <fct> NA, NA, 9/30/2015, NA, NA, NA, N...
## $ months_since_2mia <int> NA, NA, 0, NA, NA, NA, NA, NA, N...
## $ avg_mia_6m        <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 2,...
## $ max_arrear_bal_6m <int> -42, 0, 1198, -114, 0, 0, -114, ...
## $ max_mia_6m        <int> 0, 0, 2, 0, 0, 0, 0, 0, 0, 2,...
## $ avg_bal_6m        <int> 132080, 116972, 128500, 36610, 7...
## $ avg_bureau_score_6m <int> 542, 494, 290, 460, 468, 484, 51...
```

```
## $ cc_util <dbl> 0.4578, 0.6299, 0.6331, 0.4990, ...
## $ annual_income <int> 76749, 78451, 31038, 56663, 7701...
## $ emp_length <int> 3, 10, 3, 8, 10, 3, 11, 5, 4, 1,...
## $ months_since_recent_cc_delinq <int> 11, 7, 6, 6, 3, 3, 13, 5, 3, 14,...
```

Create default_event

```
oneypd_tree <- mutate(oneypd_tree, default_event = if_else(
  oneypd_tree$arrears_event == 1 |
  oneypd_tree$bankrupt_event == 1 |
  oneypd_tree$term_expiry_event == 1,
  1,0))
```

Create default flag

From “default_event” derive “default_indicator” as “Yes” “No”

```
oneypd_tree <- mutate(oneypd_tree, default_indicator = if_else(
  oneypd_tree$default_event == 1, "Yes", "No"
))
oneypd_tree$default_indicator <- as.factor(oneypd_tree$default_indicator)
```

1.2 Select a subset of variables

```
oneypd_tree_sel_orig <- dplyr::select(oneypd_tree, "default_indicator", "default_event", "bureau_score"
  "num_ccj", "time_since_ccj", "ccj_amount", "ltv", "mob", "max_arrears_12m",
  "max_arrears_bal_6m", "avg_bal_6m", "annual_income", "loan_balance", "loan_term",
  "cc_util", "emp_length", "months_since_recent_cc_delinq")
```

1.3 Filter out NAs

```
oneypd_tree_sel <- na.omit(oneypd_tree_sel_orig)
```

1.4 Split train/test

```
set.seed(123)
train_index <- createDataPartition(oneypd_tree_sel$default_event, p = 0.70, list=FALSE)
train <- oneypd_tree_sel[train_index, ]
test <- oneypd_tree_sel[-train_index, ]
```

Perform Random Forest analysis

2.1 Fit random forest

```
set.seed(123)
rf_oneypd <- randomForest(default_indicator ~ . - default_event, data = oneypd_tree_sel[train_index, ],
                           importance=TRUE, na.action=na.omit)
```

2.2 Variable importance analysis

```
imp <- importance(rf_oneypd)
print(imp)
```


	No	Yes	MeanDecreaseAccuracy
bureau_score	7.439256	12.1392643	12.319947
time_since_bankrupt	4.105219	-1.8075923	3.474521
num_ccj	3.127612	-0.8233618	3.216921
time_since_ccj	5.848557	-1.1879736	5.751842
ccj_amount	3.729015	-1.6011149	3.539973
ltv	9.960364	-1.6058195	9.377171
mob	8.026681	7.5793098	10.924976
max_arrears_12m	13.963881	16.0331846	22.442993
max_arrears_bal_6m	11.014328	12.1261976	14.149976
avg_bal_6m	12.155950	-8.5200624	11.754771
annual_income	10.577837	21.8618559	19.541862
loan_balance	11.727443	-7.9372367	11.236998
loan_term	17.160199	32.8372513	25.907522
cc_util	12.908747	59.4728637	39.458262
emp_length	7.377317	1.7508620	8.476810
months_since_recent_cc_delinq	4.821820	4.6924535	7.560639

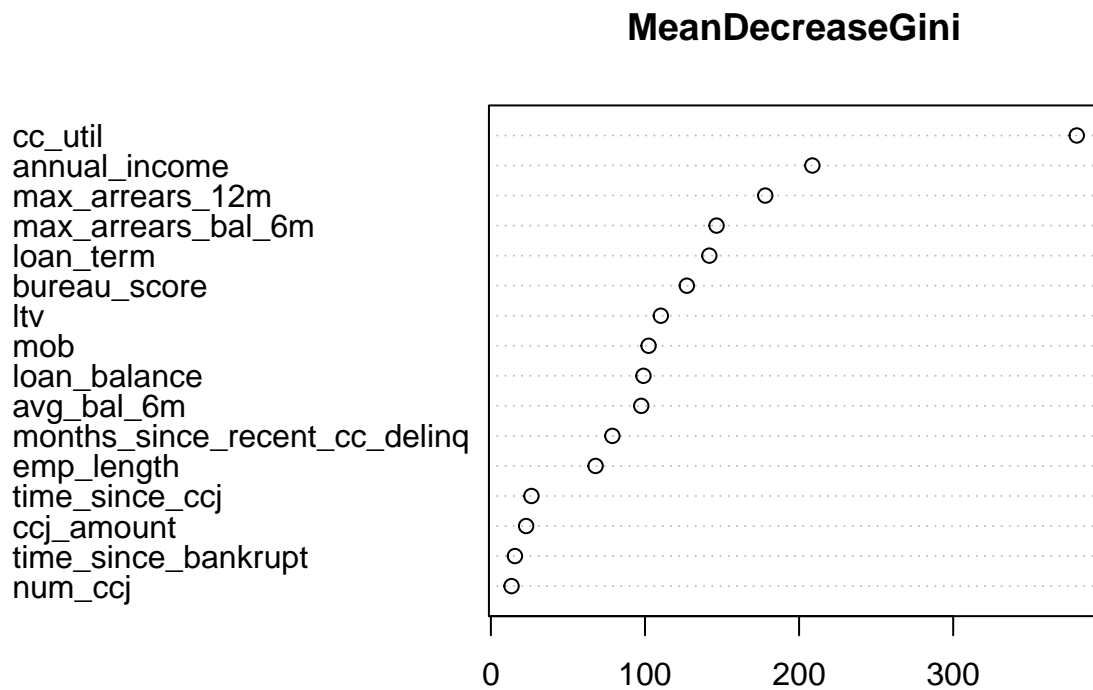
	MeanDecreaseGini
bureau_score	127.14346
time_since_bankrupt	15.60141
num_ccj	13.40092
time_since_ccj	26.31201
ccj_amount	22.87911
ltv	110.28143
mob	102.33820
max_arrears_12m	178.03480
max_arrears_bal_6m	146.42412
avg_bal_6m	97.56750
annual_income	208.59031
loan_balance	98.99793
loan_term	141.74390
cc_util	380.16297
emp_length	67.93838
months_since_recent_cc_delinq	78.84087

```
for (i in 3:4){
  ord <- order(imp[,i], decreasing=FALSE)
  dotchart(imp[ord, i], main=colnames(imp)[i], )
}
```

MeanDecreaseAccuracy


cc_util
loan_term
max_arrears_12m
annual_income
max_arrears_bal_6m
bureau_score
avg_bal_6m
loan_balance
mob
ltv
emp_length
months_since_recent_cc_delinq
time_since_ccj
ccj_amount
time_since_bankrupt
num_ccj






3. Perform boosting analysis

```
set.seed(1)
boost_oneypd = gbm(default_event ~ . - default_indicator , data = oneypd_tree_sel[train_index, ],
                    distribution='gaussian', n.trees=100, interaction.depth=4)
summary(boost_oneypd)
```



```
##                                var      rel.inf
## cc_util                      cc_util 38.18022625
## max_arrears_12m              max_arrears_12m 21.19633822
## annual_income                annual_income 13.11403783
## loan_term                    loan_term 10.76196486
## bureau_score                 bureau_score 5.63547561
## max_arrears_bal_6m           max_arrears_bal_6m 3.22845020
## mob                          mob 3.17770493
## avg_bal_6m                   avg_bal_6m 1.57542896
## ltv                          ltv 1.11793705
## months_since_recent_cc_delinq months_since_recent_cc_delinq 0.81599030
## loan_balance                 loan_balance 0.36175061
## emp_length                   emp_length 0.33470182
## ccj_amount                   ccj_amount 0.29275840
## time_since_ccj               time_since_ccj 0.11924127
## time_since_bankrupt          time_since_bankrupt 0.05333382
## num_ccj                      num_ccj 0.03465986
```

```
plot(boost_oneypd, i.var='max_arrears_12m')
```



3.1 Test sample analysis

```
yhat_boost_oneydpd = predict(boost_oneydpd, newdata = oneydpd_tree_sel[-train_index, ],
                              n.trees = 100
                              )
oneydpd_test_boost = oneydpd_tree_sel[-train_index, 'default_event']
mean((yhat_boost_oneydpd - oneydpd_test_boost)^2)
```

```
## [1] 0.02794844
```

```
summary(yhat_boost_oneydpd)
```

```
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max.
## -0.054841 -0.002370  0.005617  0.052472  0.030983  1.265853
```

3.2 Inclusion of shrinkage

```
boost_oneydpd_1 = gbm(default_event ~ . - default_indicator,
                      data=oneydpd_tree_sel[train_index, ], distribution='gaussian',
                      n.trees=100, interaction.depth=4, shrinkage=0.20, verbose=F, cv.folds=5)
yhat_oneydpd_1 = predict(boost_oneydpd_1,
```




```

newdata=oneypd_tree_sel[-train_index,],
n.trees=100)
mean((yhat_oneypd_1 - oneypd_test_boost)^2)

```

```
## [1] 0.02869406
```

```
summary(boost_oneypd_1)
```



```


##               var      rel.inf
## cc_util          cc_util 33.3340780
## max_arrears_12m    max_arrears_12m 24.3333089
## annual_income      annual_income 13.0763450
## loan_term          loan_term 10.0117294
## bureau_score       bureau_score  4.6593146
## mob                mob         3.6244708
## max_arrears_bal_6m max_arrears_bal_6m 3.6152022
## ltv                ltv         1.7703221
## avg_bal_6m         avg_bal_6m  1.6727006
## months_since_recent_cc_delinq months_since_recent_cc_delinq 1.1097567
## time_since_ccj      time_since_ccj  0.9762302
## emp_length          emp_length  0.7957802
## ccj_amount          ccj_amount  0.3722192
## loan_balance        loan_balance  0.2522918
## time_since_bankrupt time_since_bankrupt 0.2169441
## num_ccj            num_ccj     0.1793060

```

Return the Optimal Number of Iterations

```
best.iter <- gbm.perf(boost_oneypd_1, method = "OOB", plot.it = TRUE,)
```


OOB generally underestimates the optimal number of iterations although predictive performance is rea




```
print(best.iter)
```

```
## [1] 27
## attr(,"smoother")
## Call:
## loess(formula = object$oobag.improve ~ x, enp.target = min(max(4,
##   length(x)/10), 50))
##
## Number of Observations: 100
## Equivalent Number of Parameters: 8.32
## Residual Standard Error: 0.0001924622
```

```
par(mfrow=c(1,2))
plot.gbm(boost_oneypd_1, i='cc_util')
```



```
plot.gbm(boost_oneypd_1, i='max_arrears_12m')
```



ML Calibration

Create the data set and fit calibration function

```
pred_orig <- as.matrix(predict(rf_oneypd, newdata = oneypd_tree_sel, type='prob'))
rf_pred <- as.matrix(pred_orig[,2])
rf_db_cal <- as.data.frame(cbind(oneypd_tree_sel$default_event, rf_pred))
colnames(rf_db_cal) <- c('def', 'pred')
```

Fit the calibration function

```
pd_model <- glm(def ~ pred, family=binomial(link='logit'), data=rf_db_cal)
summary(pd_model)
```

```
##
## Call:
## glm(formula = def ~ pred, family = binomial(link = "logit"),
##      data = rf_db_cal)
##
## Deviance Residuals:
```


```
##      Min      1Q   Median      3Q      Max
## -3.2682 -0.0845 -0.0793 -0.0793  3.3955
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -5.7616     0.1010  -57.04  <2e-16 ***
## pred         12.9038     0.2811   45.90  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 10477.1  on 25468  degrees of freedom
## Residual deviance:  1991.2  on 25467  degrees of freedom
## AIC: 1995.2
##
## Number of Fisher Scoring iterations: 8
```

ML Model validation

```
# ROC analysis

predict_test_orig <- as.matrix(
  predict(rf_oneypd, newdata = oneypd_tree_sel[-train_index, ], type='prob'))
predict_test <- as.matrix(predict_test_orig[,2])
oneypd_test <- oneypd_tree_sel[-train_index, 'default_indicator']
actual_test <- as.matrix(ifelse(oneypd_test=='Yes', 1, 0))
pred_test <- prediction(predict_test, actual_test)
perf_test <- performance(pred_test, 'tpr', 'fpr')
auc_test <- performance(pred_test, 'auc')

plot(perf_test, main='ROC Curve', sub=paste('AUC:', round(auc_test@y.values[[1]][1],5)), colorize=T)
abline(0,1, lty=8, col='black')
```



```
#KS (Kolmogrov-Smirnov) Analysis
ks_test <- max(attr(perf_test,'y.values')[[1]] - attr(perf_test,'x.values')[[1]])
print(paste('KS Test:', round(ks_test,5)))
```

```
## [1] "KS Test: 0.73036"
```

```
gini_test <- giniCoef(predict_test, actual_test)
print(paste('Gini Index:',round(gini_test,5)))
```

```
## [1] "Gini Index: 0.8776"
```

Calibrated PD Validation

```
# predict base on model params
rf_db_cal$pd <- predict(pd_model, new_data = rf_deb_cal, type='response')

#create bands
score_cust <- smbinning.custom(rf_db_cal, y='def', x='pred',
                              cuts=c(0.2, 0.4, 0.6, 0.8))

rf_db_cal <- smbinning.gen(rf_db_cal, score_cust, chrname='band')
```

```

# Compare actual v. fitted values
# calc mean values
rf_db_cal_plot <- rf_db_cal %>%
  dplyr::group_by(band) %>%
  dplyr::summarise(mean_dr = round(mean(def), 4),
                    mean_pd = round(mean(pd), 4))

# Compute RMSe
rmse <- sqrt(mean((rf_db_cal_plot$mean_dr - rf_db_cal_plot$mean_pd)^2))
round(rmse,5)

```

```
## [1] 0.02719
```

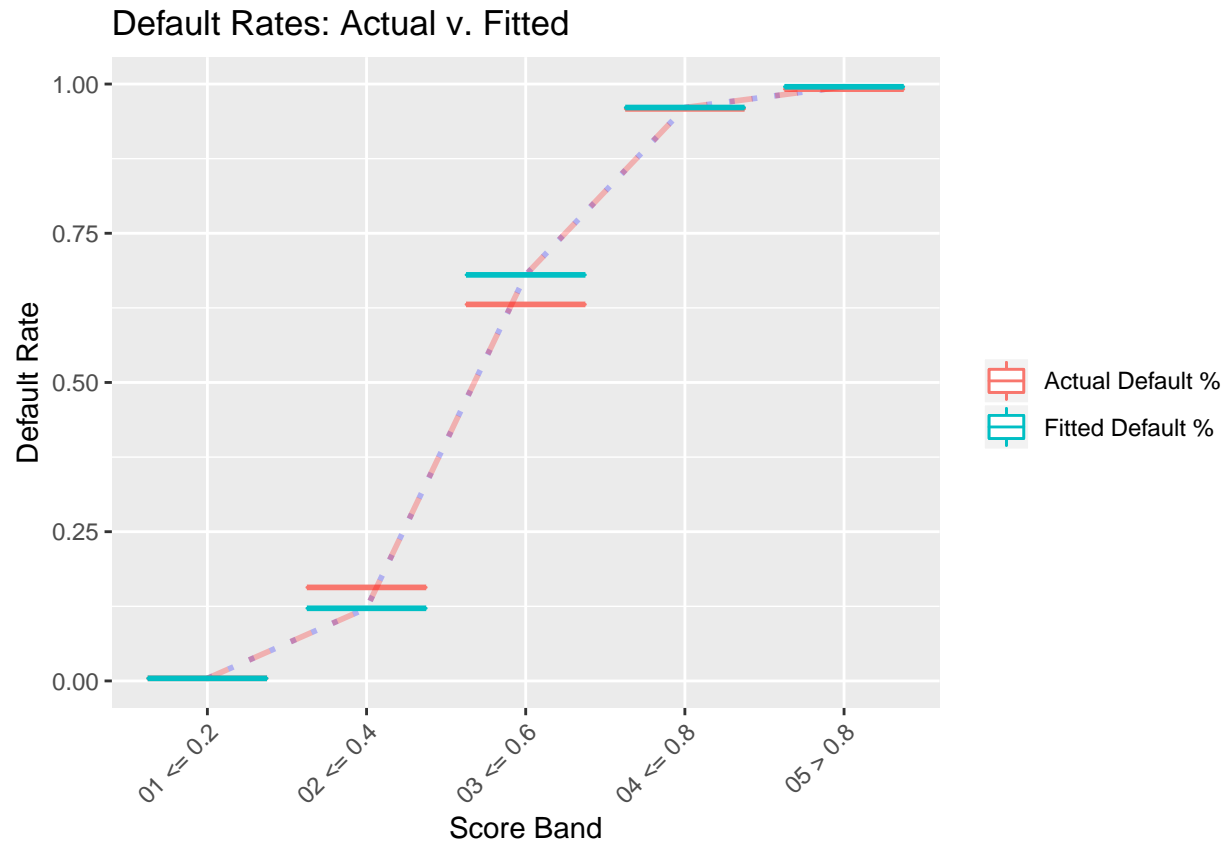
```

ggplot(data=rf_db_cal_plot, aes(x=band)) +
  geom_line(y=rf_db_cal_plot$mean_pd, group=1, color='blue', lwd=1,alpha=.25, linetype='dotted') +
  geom_boxplot(aes(y=rf_db_cal_plot$mean_dr, color='blue')) +

  geom_line(y=rf_db_cal_plot$mean_pd, group=1, color='red', lwd=1,alpha=.25, linetype='dashed') +
  geom_boxplot(aes(y=rf_db_cal_plot$mean_pd, color='red')) +

  ylab('Default Rate') +
  xlab('Score Band') +
  ggtitle("Default Rates: Actual v. Fitted") +
  scale_color_discrete(name = "", labels = c("Actual Default %", "Fitted Default %")) +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))

```



““